

---

# Convolutional Neural Networks

## Introduction

### Table of Contents

Clean .....	1
Definition .....	2
Overview .....	2
Overview .....	2
Convolution .....	3
Pooling .....	3
Convolution .....	3
CNN Load Data .....	5
CNN Architecture .....	5
CNN Training .....	6
CNN Applying .....	6
References .....	7

Augusto Damasceno

[augustodamasceno@protonmail.com](mailto:augustodamasceno@protonmail.com)

augustodamasceno.org

Copyright (c) 2025, Augusto Damasceno.

All rights reserved.

SPDX-License-Identifier: CC-BY-4.0

Images in the dataset folder are from the "Vehicle Detection Image Set" by Baris Dincer on Kaggle.

Source: <https://www.kaggle.com/datasets/brsdincer/vehicle-detection-image-set>

License: Open Data Commons Database Contents License v1.0

SPDX-License-Identifier: ODC-DbCL-1.0

## Clean

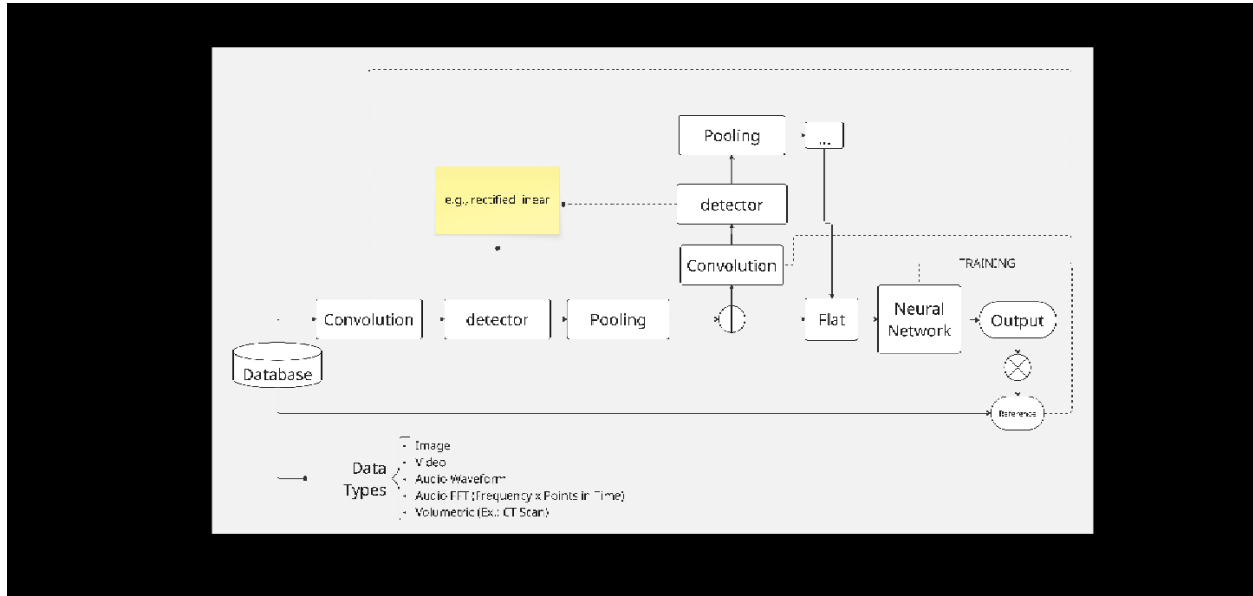
```
% Removes all variables from the workspace.
clear;
% Close all figures and
close all;
% Clear Command Window
clc
```

# Definition

Convolutional networks are simply neural networks that use convolution in place of general matrix multiplication in at least one of their layers. [1]

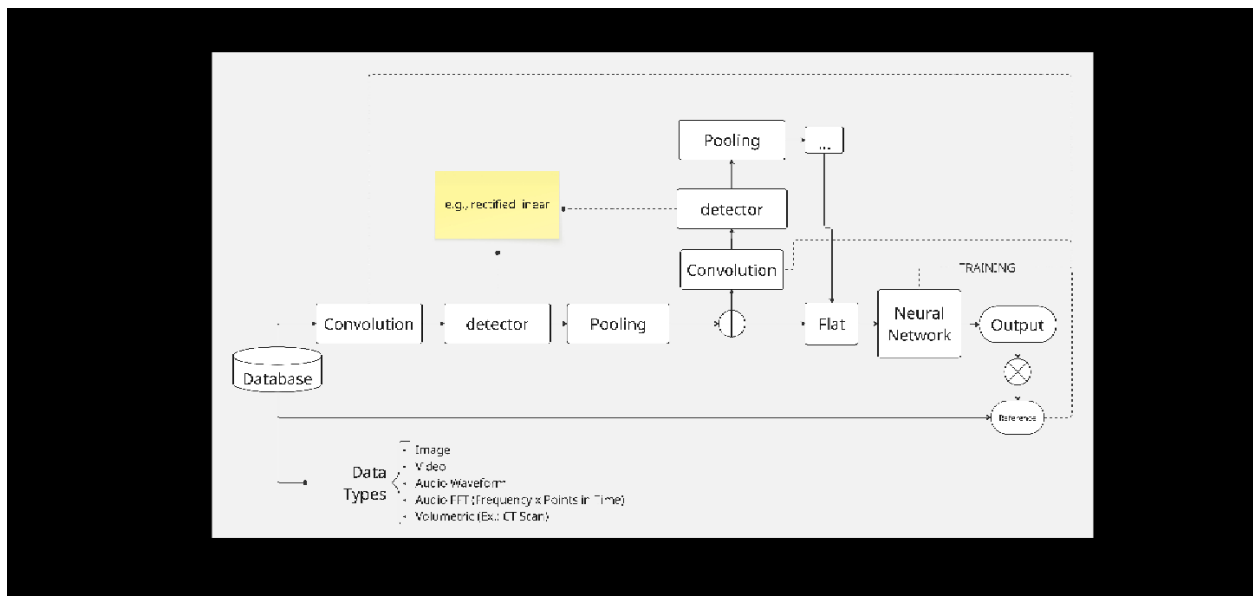
# Overview

```
imshow("overview.jpg");
```



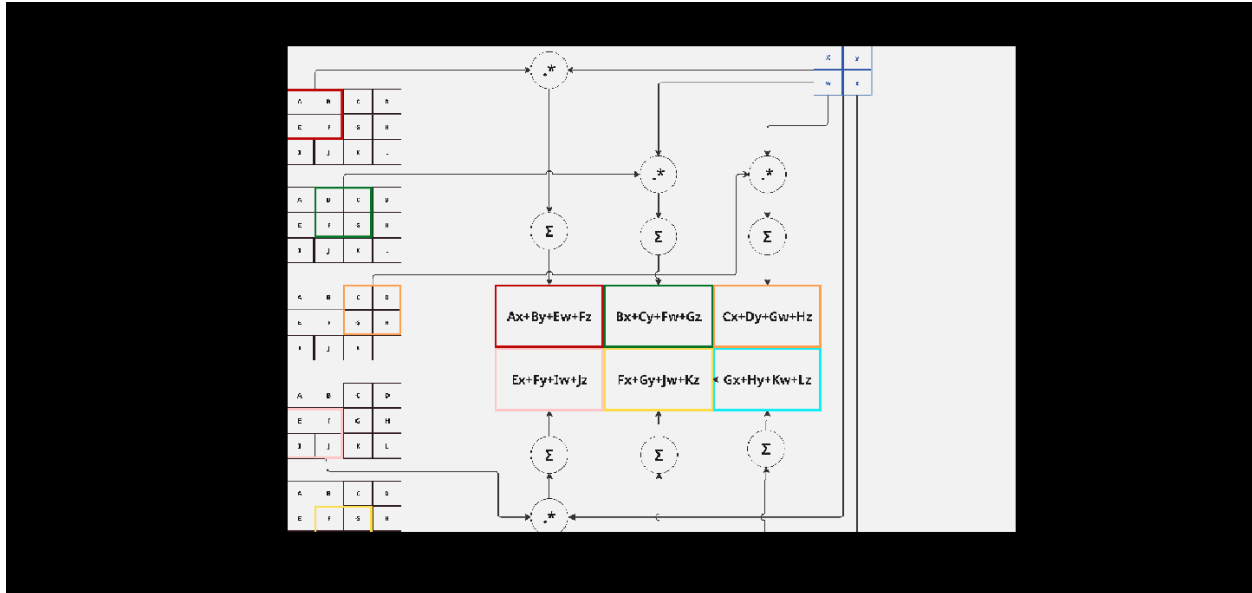
# Overview

```
imshow("overview.jpg");
```




# Convolution

```
imshow("convolution.jpg");
```



# Pooling

```
imshow("pooling.jpg");
```



# Convolution

```
kernels = {  
    [0 0 0; 0 1 0; 0 0 0]; % identity
```

```

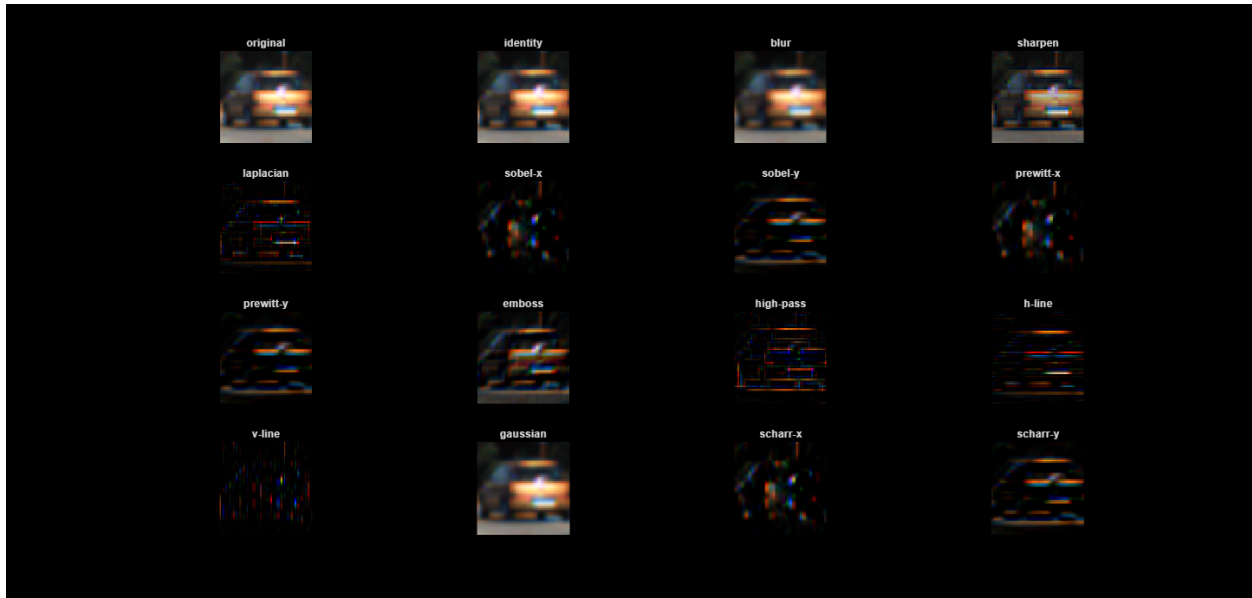
ones(3)/9;                                % box blur
[0 -1 0; -1 5 -1; 0 -1 0];               % sharpen
[0 -1 0; -1 4 -1; 0 -1 0];               % laplacian
[-1 0 1; -2 0 2; -1 0 1];                % sobel-x
[-1 -2 -1; 0 0 0; 1 2 1];                % sobel-y
[-1 0 1; -1 0 1; -1 0 1];                % prewitt-x
[-1 -1 -1; 0 0 0; 1 1 1];                % prewitt-y
[-2 -1 0; -1 1 1; 0 1 2];                % emboss
[1 1 1; 1 -8 1; 1 1 1];                  % high-pass (hp8)
[-1 -1 -1; 2 2 2; -1 -1 -1];             % horizontal line
[-1 2 -1; -1 2 -1; -1 2 -1];             % vertical line
[1 2 1; 2 4 2; 1 2 1]/16                 % gaussian
[-3 0 3; -10 0 10; -3 0 3]/16            % scharr-x
[-3 -10 -3; 0 0 0; 3 10 3]/16            % scharr-y
};

img = imread('dataset/vehicles/7.png');
img = im2double(img);
[H, W, C] = size(img);
dimStr = sprintf('Image dimensions: %d x %d x %d', H, W, C);

figure(1);
filter_names = {'identity','blur','sharpen','laplacian', ...
                'sobel-x','sobel-y','prewitt-x','prewitt-y', ...
                'emboss', 'high-pass', 'h-line', 'v-line', ...
                'gaussian', 'scharr-x', 'scharr-y'};

subplot(4,4,1), imshow(myim2uint8(img)), title('original')
if ~exist('samples','dir'), mkdir('samples'); end
imwrite(myim2uint8(img), fullfile('samples', ...
    ['original.png']));
for n = 1:numel(kernels)
    K = kernels{n};
    out = zeros(H-2,W-2, 3);
    for i = 2:H-1
        for j = 2:W-1
            p = img(i-1:i+1,j-1:j+1,:);
            out(i-1,j-1, 1) = max(sum(p(:, :, 1).*K, 'all'), 0);
            out(i-1,j-1, 2) = max(sum(p(:, :, 2).*K, 'all'), 0);
            out(i-1,j-1, 3) = max(sum(p(:, :, 3).*K, 'all'), 0);
        end
    end
    subplot(4,4,n+1), imshow(myim2uint8(out)), title(filter_names{n})
    imwrite(myim2uint8(out), fullfile('samples', ...
        ['convolution_' filter_names{n} '.png']));
end
end

```



## CNN Load Data

```
imds = imageDatastore(fullfile("dataset", {'vehicles', 'non-vehicles'}), ...
    'LabelSource', 'foldernames', 'FileExtensions', '.png');

imgCell = readall(imds);
trainLabs = imds.Labels;
trainImgs = cat(4, imgCell{:});
```

## CNN Architecture

```
layers = [
    imageInputLayer([64 64 3])
    convolution2dLayer(3, 8, "Padding", "same")
    batchNormalizationLayer
    reluLayer
    maxPooling2dLayer(2, "Stride", 2)

    convolution2dLayer(3, 16, "Padding", "same")
    batchNormalizationLayer
    reluLayer
    maxPooling2dLayer(2, "Stride", 2)

    fullyConnectedLayer(2)
    softmaxLayer
    classificationLayer];

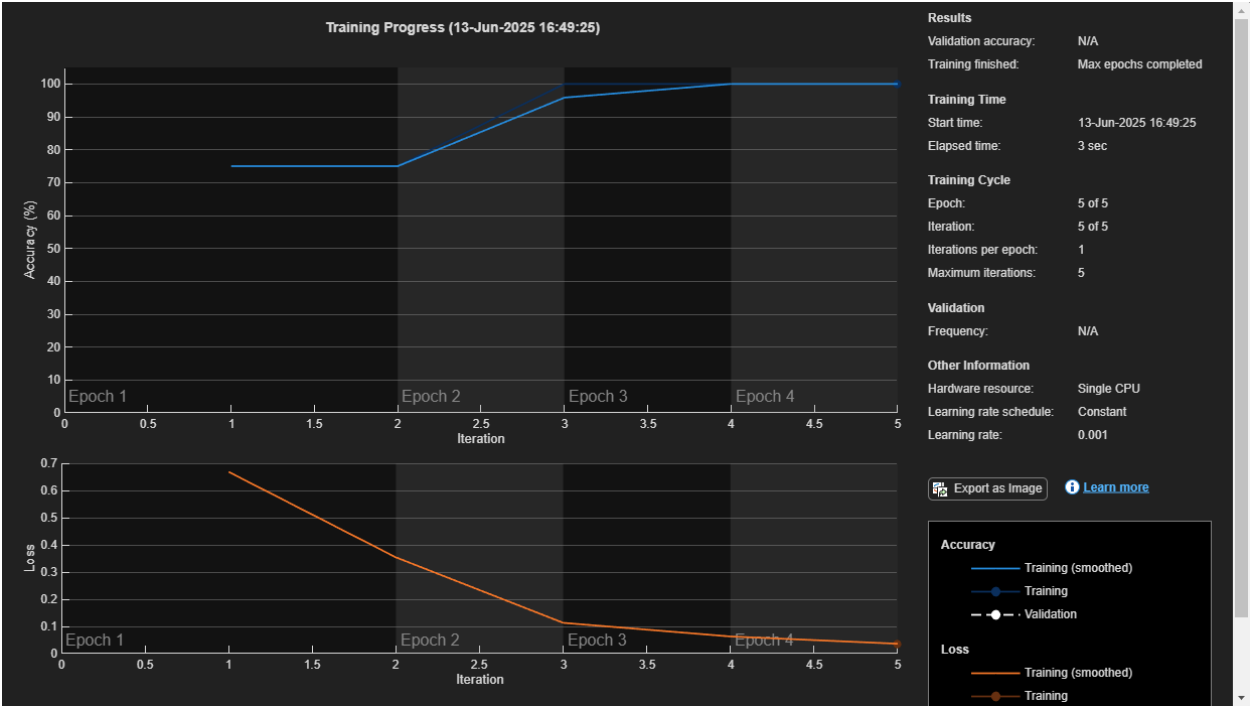
opts = trainingOptions("adam", ...
    "MaxEpochs", 5, "MiniBatchSize", 64, ...
    "Shuffle", "every-epoch", ...
    "Plots", "training-progress");
```

# CNN Training

```
net = trainNetwork(trainImgs, trainLabs, layers, opts);
```

Training on single CPU.  
Initializing input data normalization.

=====
=====
Epoch   Iteration   Time Elapsed   Mini-batch   Mini-batch
Base Learning
(hh:mm:ss)   Accuracy   Loss
Rate
=====
=====
1   1   00:00:02   75.00%   0.6688
0.0010
5   5   00:00:03   100.00%   0.0363
0.0010
=====
=====
Training finished: Max epochs completed.



# CNN Applying

```
predLabs = classify(net, imds);
```

```
tbl = table(imds.Files, imds.Labels, predLabs, ...  
           'VariableNames',{'file','true','pred'});  
writetable(tbl, 'cnn_results.csv')
```

## References

- [1] I. Goodfellow, Y. Bengio, and A. Courville, Deep Learning. Cambridge, MA: MIT Press, 2016, p. 330, para. 1. ISBN: 978-0-262-03561-3.
- [2] A. Karpathy, J. Johnson, and L. Fei-Fei, CS231n: Convolutional Neural Networks for Visual Recognition. [Online]. Available: <https://cs231n.stanford.edu/>. [Accessed: Jun 9, 2025].
- [3] G. Sanderson, "But What Is a Convolution?," **YouTube**, Nov. 18, 2022. [Online]. Available: <https://youtu.be/KuXjwB4LzSA>. [Accessed: Jun 9, 2025].
- [4] G. Sanderson, "discrete.py," **GitHub**, 2022. [Online]. Available: [https://github.com/3b1b/videos/blob/master/\\_2022/convolutions/discrete.py](https://github.com/3b1b/videos/blob/master/_2022/convolutions/discrete.py) [Accessed: Jun 9, 2025].

*Published with MATLAB® R2025a*